


# Data cleaning and preparation

## Data Science Workflow



## WHAT DATA SCIENTISTS SPEND THE MOST TIME DOING



Source: CrowdFlower 2016

- **Pandas:** provide a high-level, flexible, and fast set of tools for cleaning the data
- **Cleaning tasks:**
  1. Validating data types
  2. Inconsistency
  3. Duplicates
  4. Missing data
  5. Extreme values / Outliers

The dataset we will work on:

In [118...]

```
import pandas as pd
import numpy as np
df = pd.read_csv("data/online_retail_II 2_noisy.csv")
df
```

C:\Users\asabb\AppData\Local\Temp\ipykernel\_19756\1257127318.py:3: DtypeWarning: Columns (3,4) have mixed types. Specify dtype option on import or set low\_memory=False.  
df = pd.read\_csv("data/online\_retail\_II 2\_noisy.csv")

Out[118...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Invoice
<b>0</b>	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	6.95	83.40	13.344	96.744	12/1
<b>1</b>	489434	79323P	PINK CHERRY LIGHTS	12	6.75	81.00	12.960	93.960	12/1
<b>2</b>	489434	79323W	WHITE CHERRY LIGHTS	12	6.75	81.00	12.960	93.960	12/1
<b>3</b>	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	2.1	100.80	16.128	116.928	12/1
<b>4</b>	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	1.25	30.00	4.800	34.800	12/1
...	...	...	...	...	...	...	...	...	...
<b>805547</b>	581587	22899	CHILDREN'S APRON DOLLY GIRL	6.0	2.1	12.60	2.016	14.616	12/9
<b>805548</b>	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4.0	4.15	16.60	2.656	19.256	12/9
<b>805549</b>	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4.0	4.15	16.60	2.656	19.256	12/9
<b>805550</b>	581587	22138	BAKING SET 9 PIECE RETROSPOT	3.0	4.95	14.85	2.376	17.226	12/9
<b>805551</b>	581587	POST	POSTAGE	1.0	18.0	18.00	2.880	20.880	12/9

805552 rows × 12 columns



To do data cleaning properly, You should be familiar with:

- Business Domain
- Context how data is being collected
- Requirements
- Data dictionary

## 1. Validate data types

### Data Preparation - Validating Data Types

Before					After				
ID	Name	Age	Weight	Country	ID	Name	Age	Weight	Country
0	Sami	"27"	75	"Palestine"	0	Sami	27	75	"Palestine"
1	Sara	"30"	68	"US"	1	Sara	30	68	"US"
2	Salwa		65	"Egypt"	2	Salwa		65	"Egypt"

- Data types define the **kind of data** that can be stored in each column of your dataset
- **integers, floats, strings, and dates**
- Ensuring that each column has the correct data type is vital for **accurate analysis and functioning of algorithm**
- Determine type based on your **understanding of the dataset**
- **If discrepancies found -> convert**

In [119...]

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 805552 entries, 0 to 805551
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Invoice           805552 non-null   int64  
 1   StockCode          805526 non-null   object  
 2   Description        805526 non-null   object  
 3   Quantity           800989 non-null   object  
 4   Price              792162 non-null   object  
 5   Total               805552 non-null   float64 
 6   Tax 16%            805552 non-null   float64 
 7   Gross              805552 non-null   float64 
 8   InvoiceDate        805552 non-null   object  
 9   Payment Method     805226 non-null   object  
 10  Customer ID       804307 non-null   float64 
 11  Country            805552 non-null   object  
dtypes: float64(4), int64(1), object(7)
memory usage: 73.8+ MB

```

## Convert *Quantity* to float/int

```

In [120... pattern = r'\d+(\.\d+)?$'

df2=df[(df['Quantity'].notna()) & (~df['Quantity'].astype(str).str.match(pattern))]
df2

```

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	InvoiceDate
<b>13</b>	489436	21755	LOVE BUILDING BLOCK WORD	"18	5.45	98.10	15.6960	113.7960	12/1/2009:0
<b>15</b>	489436	84879	ASSORTED COLOUR BIRD ORNAMENT	"16	1.69	27.04	4.3264	31.3664	12/1/2009:0
<b>34</b>	489437	21364	PEACE SMALL WOOD LETTERS	"2	6.75	13.50	2.1600	15.6600	12/1/2009:0
<b>48</b>	489437	22271	FELTCRAFT DOLL ROSIE	"6	2.95	17.70	2.8320	20.5320	12/1/2009:0
<b>447</b>	489522	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	"1	3.75	3.75	0.6000	4.3500	12/1/2009:0



```
In [121... df['Quantity'] = df['Quantity'].astype(str).str.replace("'", ' ')
df['Quantity'] = df['Quantity'].astype(float)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 805552 entries, 0 to 805551
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Invoice          805552 non-null   int64  
 1   StockCode         805526 non-null   object  
 2   Description       805526 non-null   object  
 3   Quantity          800989 non-null   float64 
 4   Price             792162 non-null   object  
 5   Total              805552 non-null   float64 
 6   Tax 16%           805552 non-null   float64 
 7   Gross             805552 non-null   float64 
 8   InvoiceDate       805552 non-null   object  
 9   Payment Method    805226 non-null   object  
 10  Customer ID      804307 non-null   float64 
 11  Country           805552 non-null   object  
dtypes: float64(5), int64(1), object(6)
memory usage: 73.8+ MB
```

## Convert *Price* to float

```
In [122... pattern = r'\d+(\.\d+)?$'
df[(df['Price'].notna()) & (~df['Price'].astype(str).str.match(pattern))]
```

Out[122...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	InvoiceDate
<b>262288</b>	522622	20966	SANDWICH BATH SPONGE	10.0	1.25	12.5	2.000	14.500	9/15/2010 15:32
<b>490597</b>	547066	21080	SET/20 RED RETROSPOT PAPER NAPKINS	12.0	0.85\$	10.2	1.632	11.832	3/20/2011 13:41



```
In [123... df['Price'] = df['Price'].astype(str).str.replace(',', '.', regex=False).str.replace(' ', '')
df['Price'] = df['Price'].astype(float)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 805552 entries, 0 to 805551
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Invoice           805552 non-null   int64  
 1   StockCode          805526 non-null   object  
 2   Description        805526 non-null   object  
 3   Quantity           800989 non-null   float64 
 4   Price              792162 non-null   float64 
 5   Total               805552 non-null   float64 
 6   Tax 16%            805552 non-null   float64 
 7   Gross              805552 non-null   float64 
 8   InvoiceDate        805552 non-null   object  
 9   Payment Method     805226 non-null   object  
 10  Customer ID       804307 non-null   float64 
 11  Country            805552 non-null   object  
dtypes: float64(6), int64(1), object(5)
memory usage: 73.8+ MB
```

## convert *customer type* to string

```
In [124...]: ## Customer
df['Customer ID'] = df['Customer ID'].fillna(-1)
df['Customer ID'] = df['Customer ID'].astype(int).astype(str)
df['Customer ID'] = df['Customer ID'].replace("-1", np.nan)
df['Customer ID'].info()

<class 'pandas.core.series.Series'>
RangeIndex: 805552 entries, 0 to 805551
Series name: Customer ID
Non-Null Count  Dtype  
----- 
804307 non-null  object  
dtypes: object(1)
memory usage: 6.1+ MB
```

## Convert *Invoice Date* to date time

```
In [125...]: df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'], format='%m/%d/%Y %H:%M')
df
```

Out[125...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Invoic...
<b>0</b>	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12.0	6.95	83.40	13.344	96.744	2009-07
<b>1</b>	489434	79323P	PINK CHERRY LIGHTS	12.0	6.75	81.00	12.960	93.960	2009-07
<b>2</b>	489434	79323W	WHITE CHERRY LIGHTS	12.0	6.75	81.00	12.960	93.960	2009-07
<b>3</b>	489434	22041	RECORD FRAME 7" SINGLE SIZE	48.0	2.10	100.80	16.128	116.928	2009-07
<b>4</b>	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24.0	1.25	30.00	4.800	34.800	2009-07
...	...	...	...	...	...	...	...	...	...
<b>805547</b>	581587	22899	CHILDREN'S APRON DOLLY GIRL	6.0	2.10	12.60	2.016	14.616	2011-12
<b>805548</b>	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4.0	4.15	16.60	2.656	19.256	2011-12
<b>805549</b>	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4.0	4.15	16.60	2.656	19.256	2011-12
<b>805550</b>	581587	22138	BAKING SET 9 PIECE RETROSPOT	3.0	4.95	14.85	2.376	17.226	2011-12
<b>805551</b>	581587	POST	POSTAGE	1.0	18.00	18.00	2.880	20.880	2011-12

805552 rows × 12 columns



## 2. Data inconsistency

# Data Preparation - Inconsistency

- Spelling
- Unit differences
- Inconsistent categorization

Before					After				
ID	Name	Age	Weight	Country	ID	Name	Age	Weight	Country
0	Sami	"27"	75	"Palestine"	0	Sami	27	75	"Palestine"
1	Sara	"30"	150	"US"	1	Sara	30	68	"US"
2	Salwa		65	"Egypt"	2	Salwa		65	"Egypt"

Data inconsistencies occur when similar data is recorded in **different formats** or **representations**, leading to **unreliable analysis**

The popular approach in finding inconsistency is using **value\_counts()** function

## Find inconsistencies in *Payment Method*

```
In [126...]: df['Payment Method'].value_counts()
```

```
Out[126...]: Payment Method
Credit Card    804002
Cash           1193
credit card     23
CC              7
cash             1
Name: count, dtype: int64
```

```
In [127...]: df['Payment Method'] = df['Payment Method'].replace({"CC": "Credit Card", "credit c"})
df['Payment Method'].value_counts()
```

```
Out[127...]: Payment Method
Credit Card    804032
Cash           1194
Name: count, dtype: int64
```

## Find inconsistencies in *Country*

```
In [128...]: df['Country'].value_counts()
```

```
Out[128...]:
```

Country	
United Kingdom	725242
Germany	16694
EIRE	15743
France	13813
Netherlands	5088
Spain	3719
Belgium	3068
Switzerland	3011
Portugal	2446
Australia	1810
Channel Islands	1569
Italy	1468
Norway	1436
Sweden	1319
Cyprus	1155
Finland	1032
Austria	922
Denmark	798
Greece	657
Unspecified	521
Poland	512
Japan	485
United Arab Emirates	383
USA	373
Singapore	339
Israel	322
Malta	282
Iceland	253
Canada	228
Lithuania	189
RSA	122
Brazil	94
Thailand	76
European Community	60
Bahrain	59
West Indies	54
Korea	53
Lebanon	45
United States	36
Nigeria	30
Czech Republic	25
UK	10
Saudi Arabia	9
AUS	2

```
Name: count, dtype: int64
```

```
In [129...]:
```

```
df['Country'] = df['Country'].replace({'UK': "United Kingdom", "United States": "US"}  
df['Country'].value_counts()
```

```
Out[129...]: Country
United Kingdom    725252
Germany          16694
EIRE              15743
France            13813
Netherlands       5088
Spain              3719
Belgium           3068
Switzerland        3011
Portugal           2446
Australia          1812
Channel Islands   1569
Italy               1468
Norway             1436
Sweden             1319
Cyprus              1155
Finland            1032
Austria             922
Denmark             798
Greece              657
Unspecified         521
Poland              512
Japan                485
USA                  409
United Arab Emirates 383
Singapore           339
Israel              322
Malta                282
Iceland              253
Canada              228
Lithuania            189
RSA                  122
Brazil                94
Thailand              76
European Community    60
Bahrain              59
West Indies            54
Korea                  53
Lebanon                45
Nigeria                30
Czech Republic         25
Saudi Arabia            9
Name: count, dtype: int64
```

### 3. Duplicates

# Data Preparation - Handling Duplicates

Before					After				
ID	Name	Age	Weight	Country	ID	Name	Age	Weight	Country
0	Sami	"27"	75	"Palestine"	0	Sami	27	75	"Palestine"
1	Sara	"30"	68	"US"	1	Sara	30	68	"US"
2	Salwa		65	"Egypt"	2	Salwa		65	"Egypt"
0	Sami	"27"	75	"Palestine"					

- Duplicates can lead to **incorrect analysis**
- Duplicates arise:
  - Repeated data entry
  - Mergeing datasets
  - Data Collection errors

## Check for duplicates

```
In [130...]: df[df.duplicated()]
```

Out[130...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	InvoiceD
342	489517	21912	VINTAGE SNAKES & LADDERS	1.0	3.75	3.75	0.6000	4.3500	2009-12 11:34
354	489517	22130	PARTY CONE CHRISTMAS DECORATION	6.0	0.85	5.10	0.8160	5.9160	2009-12 11:34
355	489517	22319	HAIRCLIPS FORTIES FABRIC ASSORTED	12.0	0.65	7.80	1.2480	9.0480	2009-12 11:34
356	489517	21913	VINTAGE SEASIDE JIGSAW PUZZLES	1.0	3.75	3.75	0.6000	4.3500	2009-12 11:34
357	489517	21821	GLITTER STAR GARLAND WITH BELLS	1.0	3.75	3.75	0.6000	4.3500	2009-12 11:34
...	...	...	...	...	...	...	...	...	...
805320	581538	22068	BLACK PIRATE TREASURE CHEST	1.0	0.39	0.39	0.0624	0.4524	2011-12 11:34
805334	581538	23318	BOX OF 6 MINI VINTAGE CRACKERS	1.0	2.49	2.49	0.3984	2.8884	2011-12 11:34
805337	581538	22992	REVOLVER WOODEN RULER	1.0	1.95	1.95	0.3120	2.2620	2011-12 11:34
805344	581538	22694	WICKER STAR	1.0	2.10	2.10	0.3360	2.4360	2011-12 11:34
805346	581538	23343	JUMBO BAG VINTAGE CHRISTMAS	1.0	2.08	2.08	0.3328	2.4128	2011-12 11:34

25108 rows × 12 columns

**Drop duplicates**

In [131...]

```
df = df.drop_duplicates()
df
```

Out[131...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Invoiced
<b>0</b>	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12.0	6.95	83.40	13.344	96.744	2009-07
<b>1</b>	489434	79323P	PINK CHERRY LIGHTS	12.0	6.75	81.00	12.960	93.960	2009-07
<b>2</b>	489434	79323W	WHITE CHERRY LIGHTS	12.0	6.75	81.00	12.960	93.960	2009-07
<b>3</b>	489434	22041	RECORD FRAME 7" SINGLE SIZE	48.0	2.10	100.80	16.128	116.928	2009-07
<b>4</b>	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24.0	1.25	30.00	4.800	34.800	2009-07
...	...	...	...	...	...	...	...	...	...
<b>805547</b>	581587	22899	CHILDREN'S APRON DOLLY GIRL	6.0	2.10	12.60	2.016	14.616	2011-12
<b>805548</b>	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4.0	4.15	16.60	2.656	19.256	2011-12
<b>805549</b>	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4.0	4.15	16.60	2.656	19.256	2011-12
<b>805550</b>	581587	22138	BAKING SET 9 PIECE RETROSPOT	3.0	4.95	14.85	2.376	17.226	2011-12
<b>805551</b>	581587	POST	POSTAGE	1.0	18.00	18.00	2.880	20.880	2011-12

780444 rows × 12 columns



## Missing Values

# Data Preparation - Missing Values

ID	Name	Age	Weight	Country
0	Sami	"27"	75	"Palestine"
1	Sara	"30"	68	"US"
2	Salwa		65	"Egypt"

## Reasons:

- Data entry
- ETL error
- Valid missing

## Solutions:

- Re-ETL
- Impute from similar
- Drop
- Keep it missing

## Check missing values

```
In [132...]: df.isna().sum()
```

```
Out[132...]: Invoice          0
StockCode         26
Description       26
Quantity        4562
Price           13386
Total            0
Tax 16%          0
Gross            0
InvoiceDate      0
Payment Method   326
Customer ID     1245
Country          0
dtype: int64
```

## Process for handling missing values:

1. **Understand missing data:** revisit the data source and understand why the data is missing
2. **Calculate from the data itself**

### 3. Impute

- From similar data points
- constant value (mean, most frequent, 0,...)
- bfill, ffill (time series data)
- ML models (knn)

### 4. Drop

## Example

In [133...]

```
import numpy as np
import pandas as pd

# Create dataframe
import numpy as np
data = pd.DataFrame(np.random.standard_normal((7, 3)))
data.iloc[2:3, 1] = np.nan
data.iloc[4:5, 1] = np.nan
data.iloc[4:5, 2] = np.nan
data.iloc[0,1]= np.nan
data.iloc[-1,2]= np.nan
data
```

Out[133...]

	0	1	2
0	0.071148	NaN	-0.690006
1	-0.740496	-2.121285	0.870374
2	-0.563426	NaN	-0.598694
3	-0.381680	1.630362	0.908357
4	1.015674	NaN	NaN
5	1.664753	0.029204	-0.509784
6	-0.752856	0.869271	NaN

In [134...]

```
# Fill with constant value
data.fillna(0)
```

```
Out[134...]
```

	0	1	2
0	0.071148	0.000000	-0.690006
1	-0.740496	-2.121285	0.870374
2	-0.563426	0.000000	-0.598694
3	-0.381680	1.630362	0.908357
4	1.015674	0.000000	0.000000
5	1.664753	0.029204	-0.509784
6	-0.752856	0.869271	0.000000

```
In [135...]
```

```
# Fill with mean/median/max/min...
data.fillna(data.mean())
```

```
Out[135...]
```

	0	1	2
0	0.071148	0.101888	-0.690006
1	-0.740496	-2.121285	0.870374
2	-0.563426	0.101888	-0.598694
3	-0.381680	1.630362	0.908357
4	1.015674	0.101888	-0.003951
5	1.664753	0.029204	-0.509784
6	-0.752856	0.869271	-0.003951

```
In [136...]
```

```
# 'forward fill': propagate last valid observation forward
data.ffill()
```

```
Out[136...]
```

	0	1	2
0	0.071148	NaN	-0.690006
1	-0.740496	-2.121285	0.870374
2	-0.563426	-2.121285	-0.598694
3	-0.381680	1.630362	0.908357
4	1.015674	1.630362	0.908357
5	1.664753	0.029204	-0.509784
6	-0.752856	0.869271	-0.509784

```
In [137...]
```

```
# backward fill
data.bfill()
```

```
Out[137...]
```

	0	1	2
0	0.071148	-2.121285	-0.690006
1	-0.740496	-2.121285	0.870374
2	-0.563426	1.630362	-0.598694
3	-0.381680	1.630362	0.908357
4	1.015674	0.029204	-0.509784
5	1.664753	0.029204	-0.509784
6	-0.752856	0.869271	NaN

```
In [138...]
```

```
# Drop all rows that have any missing value  
data.dropna()
```

```
Out[138...]
```

	0	1	2
1	-0.740496	-2.121285	0.870374
3	-0.381680	1.630362	0.908357
5	1.664753	0.029204	-0.509784

## Back to Retail dataset

```
In [139...]
```

```
df.isna().sum()
```

```
Out[139...]
```


```
Invoice          0  
StockCode       26  
Description     26  
Quantity        4562  
Price           13386  
Total            0  
Tax 16%          0  
Gross            0  
InvoiceDate      0  
Payment Method   326  
Customer ID     1245  
Country          0  
dtype: int64
```

```
In [140...]
```

```
#!pip install missingno  
import missingno as msno  
msno.matrix(df.sort_values(by="Payment Method"))  
#print(df[df['Quantity'].isna()][['Payment Method']].value_counts(normalize=True))  
#df[df['Quantity'].notna()][['Payment Method']].value_counts(normalize=True)
```

```
Out[140...]
```

```
<Axes: >
```



## Handle Price

```
In [141]: df[df['Price'].isna()]
```

Out[141...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Invoice Date
11	489435	22353	LUNCHBOX WITH CUTLERY FAIRY CAKES	NaN	NaN	30.60	4.8960	35.4960	2009-10-07:2
47	489437	20971	PINK BLUE FELT CRAFT TRINKET BOX	NaN	NaN	15.00	2.4000	17.4000	2009-10-09:0
52	489437	22111	SCOTTIE DOG HOT WATER BOTTLE	NaN	NaN	14.85	2.3760	17.2260	2009-10-09:0
60	489438	21411	GINGHAM HEART DOORSTOP RED	32.0	NaN	80.00	12.8000	92.8000	2009-10-09:2
77	489439	16161P	WRAP ENGLISH ROSE	25.0	NaN	10.50	1.6800	12.1800	2009-10-09:2
...	...	...	...	...	...	...	...	...	...
805154	581496	22112	CHOCOLATE HOT WATER BOTTLE	NaN	NaN	29.70	4.7520	34.4520	2011-10-10:2
805165	581496	22190	LOCAL CAFE MUG	24.0	NaN	9.36	1.4976	10.8576	2011-10-10:2
805192	581501	21564	PINK HEART SHAPE LOVE BUCKET	24.0	NaN	18.96	3.0336	21.9936	2011-10-10:2
805349	581567	22464	HANGING METAL HEART LANTERN	24.0	NaN	18.96	3.0336	21.9936	2011-10-11:5
805517	581585	84879	ASSORTED COLOUR BIRD ORNAMENT	16.0	NaN	27.04	4.3264	31.3664	2011-10-12:3

13386 rows × 12 columns



In [142...]

```
print(df['Price'].isna().sum())
df['Price'] = df['Price'].fillna(df['Total'] / df['Quantity'])
```

```
df['Price'].isna().sum()
```

13386

```
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\4184607728.py:2: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
```

```
df['Price'] = df['Price'].fillna(df['Total'] / df['Quantity'])
```

Out[142... np.int64(4562)

In [143... *### Get help from StockCode, assuming the product usually sold with the same price*  
df = df[df['StockCode'].notna()]

```
# Step 2: Keep only numeric StockCodes  
df = df[df['StockCode'].astype(str).str.isnumeric()]
```

```
# Step 3: Convert StockCode to integer (optional)  
df["StockCode"] = df["StockCode"].astype(int)
```

```
# Step 4: Fill missing prices using the average price per StockCode  
df['Price'] = df.groupby('StockCode')['Price'].transform(lambda x: x.fillna(x.mean()))
```

In [144... df[df['Price'].isna()]

Out[144... 

Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	InvoiceDate	Payme Meth
---------	-----------	-------------	----------	-------	-------	------------	-------	-------------	---------------



## Handle Quantity

In [145... df[df['Quantity'].isna()]

Out[145...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Inv
11	489435	22353	LUNCHBOX WITH CUTLERY FAIRY CAKES	NaN	2.533019	30.60	4.8960	35.4960	20
47	489437	20971	PINK BLUE FELT CRAFT TRINKET BOX	NaN	1.253953	15.00	2.4000	17.4000	20
52	489437	22111	SCOTTIE DOG HOT WATER BOTTLE	NaN	4.894426	14.85	2.3760	17.2260	20
441	489522	21479	WHITE SKULL HOT WATER BOTTLE	NaN	3.877635	7.50	1.2000	8.7000	20
598	489529	20657	TROPICAL LUGGAGE TAG	NaN	1.241286	1.25	0.2000	1.4500	20
...	...	...	...	...	...	...	...	...	...
804775	581467	84976	RECTANGULAR SHAPED MIRROR	NaN	1.229000	2.37	0.3792	2.7492	20
804802	581469	21158	MOODY GIRL DOOR HANGER	NaN	0.956357	0.39	0.0624	0.4524	20
804908	581473	20718	RED RETROSPOT SHOPPER BAG	NaN	1.249688	1.25	0.2000	1.4500	20
805046	581479	22087	PAPER BUNTING WHITE LACE	NaN	2.892199	29.50	4.7200	34.2200	20
805154	581496	22112	CHOCOLATE HOT WATER BOTTLE	NaN	4.879060	29.70	4.7520	34.4520	20

4002 rows × 12 columns



In [146...]

```
print(df['Quantity'].isna().sum())
df['Quantity'] = df['Quantity'].fillna(df['Total'] / df['Price'])
df['Quantity'].isna().sum()
```

4002

Out[146...]

np.int64(0)

## Handle payment method

```
In [147... df[df['Payment Method'].isna()]  
# df.isna().sum()
```

Out[147...]

		Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	Invoice Date
				CLEAR CRYSTAL STAR PHONE CHARM	24.0	0.85	20.40	3.2640	23.6640	2009-12-15
1171	489560	90093		ECONOMY PASSPORT COVER	1.0	2.10	2.10	0.3360	2.4360	2009-12-01
12036	490685	48197		DOOR MAT BIRD ON THE WIRE	1.0	6.75	6.75	1.0800	7.8300	2009-12-13
14512	491009	22353		LUNCHBOX WITH CUTLERY FAIRY CAKES	6.0	2.55	15.30	2.4480	17.7480	2009-12-18
19280	491645	20764		ABSTRACT CIRCLES SKETCHBOOK	1.0	3.75	3.75	0.6000	4.3500	2009-12-16
	...	...	...	...	...	...	...	...	...	...
792635	580399	23355		HOT WATER BOTTLE KEEP CALM	2.0	4.95	9.90	1.5840	11.4840	2011-11-12
801801	581166	22722		SET OF 6 SPICE TINS PANTRY DESIGN	1.0	3.95	3.95	0.6320	4.5820	2011-11-14
802115	581181	22144		CHRISTMAS CRAFT LITTLE FRIENDS	6.0	2.10	12.60	2.0160	14.6160	2011-11-15
802550	581230	22617		BAKING SET SPACEBOY DESIGN	3.0	4.95	14.85	2.3760	17.2260	2011-11-10
804537	581443	22758		LARGE PURPLE BABUSHKA NOTEBOOK	12.0	0.39	4.68	0.7488	5.4288	2011-11-16

288 rows × 12 columns



```
In [148...]: #This fills missing payment methods with the first non-null value from the same Invoice
print(df['Payment Method'].isna().sum())

def find_payment_method(group):
    rows = group.dropna()
    v = rows.iloc[0] if rows.shape[0] > 0 else np.nan
    return group.fillna(v).astype(object)

df['Payment Method'] = df['Payment Method'].groupby(df['Invoice']).transform(find_payment_method)

df['Payment Method'].isna().sum()
```

288

```
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\1119359148.py:7: FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will change in a future version. Call result.infer_objects(copy=False) instead. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    return group.fillna(v).astype(object)
```

Out[148...]: np.int64(1)

## Handle customer ID

```
In [149...]: df[df['Customer ID'].isna()]
```

Out[149...]

	Invoice	StockCode	Description	Quantity	Price	Total	Tax 16%	Gross	In
<b>1951</b>	489640	84006	MAGIC TREE - PAPER FLOWERS	72.0	0.85	61.20	9.7920	70.9920	2
<b>2187</b>	489667	22246	GARLAND, MAGIC GARDEN 1.8M	12.0	1.95	23.40	3.7440	27.1440	2
<b>2372</b>	489683	48116	DOOR MAT MULTICOLOUR STRIPE	2.0	6.75	13.50	2.1600	15.6600	2
<b>2483</b>	489702	22086	PAPER CHAIN KIT 50'S CHRISTMAS	280.0	2.55	714.00	114.2400	828.2400	2
<b>2691</b>	489780	22083	PAPER CHAIN KIT RETRO SPOT	12.0	2.95	35.40	5.6640	41.0640	2
...	...	...	...	...	...	...	...	...	...
<b>802494</b>	581220	20992	JAZZ HEARTS PURSE NOTEBOOK	24.0	0.39	9.36	1.4976	10.8576	2
<b>802538</b>	581225	23341	PINK DINER WALL CLOCK	2.0	8.50	17.00	2.7200	19.7200	2
<b>803284</b>	581376	22767	TRIPLE PHOTO FRAME CORNICE	1.0	9.95	9.95	1.5920	11.5420	2
<b>803707</b>	581405	22940	FELTCRAFT CHRISTMAS FAIRY	1.0	4.25	4.25	0.6800	4.9300	2
<b>805385</b>	581571	21169	YOU'RE CONFUSING ME METAL SIGN	1.0	1.69	1.69	0.2704	1.9604	2

1109 rows × 12 columns



In [150...]


```
print(df['Customer ID'].isna().sum())
df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna(x['Customer ID'].mode()))
df['Customer ID'].isna().sum()
df.isna().sum()
```

1109


```
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: FutureWarning: Do
wncasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will chan
ge in a future version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
C:\Users\asabb\AppData\Local\Temp\ipykernel_19756\3373735653.py:2: DeprecationWarnin
g: DataFrameGroupBy.apply operated on the grouping columns. This behavior is depreca
ted, and in a future version of pandas the grouping columns will be excluded from th
e operation. Either pass `include_groups=False` to exclude the groupings or explicit
ly select the grouping columns after groupby to silence this warning.
    df['Customer ID'] = df.groupby('Invoice').apply(lambda x: x['Customer ID'].fillna
(x['Customer ID'].values[0]).astype(object)).reset_index('Invoice', drop=True)
```

```
Out[150]: Invoice      0  
StockCode     0  
Description   0  
Quantity      0  
Price         0  
Total         0  
Tax 16%       0  
Gross         0  
InvoiceDate   0  
Payment Method 1  
Customer ID   55  
Country        0  
dtype: int64
```


## Outliers




The Empirical Rule



Approximately 68% of the area under a normal curve is between one standard deviation above and below the mean.




Approximately 95% of the area under a normal curve is between two standard deviations above and below the mean.



More than 99% of the area under a normal curve is between three standard deviations above and below the mean.

## Detecting Outliers with z-Scores



Recap

$$Z = \frac{X - \mu}{\sigma}$$

Find extreme/outlier prices

```
In [151]: df['price_zscore'] = (df['Price'] - df['Price'].mean()) / df['Price'].std()
df['outlier'] = df['price_zscore'].apply(lambda x: abs(x) >= 3)
df[df['outlier']] == True].sort_values(by="price_zscore", ascending=False)
```

Out[151...]

	<b>Invoice</b>	<b>StockCode</b>	<b>Description</b>	<b>Quantity</b>	<b>Price</b>	<b>Total</b>	<b>Tax 16%</b>
			PICNIC BASKET WICKER 60 PIECES	1.000000	649.500000	649.50	103.9200
<b>563073</b>	556446	22502					75
			PICNIC BASKET WICKER 60 PIECES	60.000000	649.500000	38970.00	6235.2000
<b>563063</b>	556444	22502					4520
			VINTAGE BLUE KITCHEN CABINET	1.000000	295.000000	295.00	47.2000
<b>215781</b>	516913	22656					34
			VINTAGE RED KITCHEN CABINET	1.000000	295.000000	295.00	47.2000
<b>274680</b>	523946	22655					34
			VINTAGE BLUE KITCHEN CABINET	1.000000	295.000000	295.00	47.2000
<b>210051</b>	516164	22656					34
...	...	...	...	...	...	...	...
			UNION STRIPE WITH FRINGE HAMMOCK	1.000000	16.630000	16.63	2.6608
<b>782842</b>	579196	21922					1
			BREAD BIN DINER STYLE IVORY	2.056394	16.485169	33.90	5.4240
<b>385679</b>	535576	22847					3
			BREAD BIN DINER STYLE IVORY	2.056394	16.485169	33.90	5.4240
<b>569496</b>	557222	22847					3
			BREAD BIN DINER STYLE RED	2.072802	16.354677	33.90	5.4240
<b>783671</b>	579283	22846					3
			BREAD BIN DINER STYLE PINK	11.070295	16.205530	179.40	28.7040
<b>553272</b>	555164	22848					20

3846 rows × 14 columns

